

```

# =====
# Hilios v18 - Dark Sector Physics
# =====
# =====
# === PART 1: IMPORTS & PARAMETERS ===
# =====

import numpy as np
import cupy as cp
import time
import h5py
import os
import matplotlib.pyplot as plt
import matplotlib.gridspec as gridspec
from IPython.display import display, clear_output

sim_params = {
    # Simulation Control
    "NUM_STEPS": 50000,
    "VIS_INTERVAL": 100,
    "DT": 1e-5,
    "GRID_SIZE": 401,

    # RDU Physics Constants
    "G_RDU": 0.05,
    "K_S": 1.0,
    "C_ACCRETION_DRAG": 0.1,
    "ACCRETION_RADIUS": 15.0,
    "C_DYN_FRICTION": 0.01,

    # Knochelman Mechanism Constants
    "A_CONST": 1.2,
    "N_SENSITIVITY": 0.5,
    "C_UNIFY": 1.5e-5,
    "K_ROT": 3.0,
    "C_BALANCE": 423.0,
    "TACHYON_MASS": 0.001,

    # Supernova & Magnetar Parameters
    "SUPERNOVA_SHELL_PARTICLES": 200,
    "SUPERNOVA_EJECTION_VELOCITY": 0.5,
    "REMNANT_MASS_FRACTION": 0.2,
    "MAGNETAR_STRENGTH": 500.0,
    "GAS_CHARGE": 0.05,

```

```

# Dark Strong Force Parameters
"DARK_FORCE_STRENGTH": 0.01, # Epsilon_DS: The overall strength
"REPULSION_RADIUS": 2.5, # r_0: The distance where the force turns repulsive

# Particle Population
"gas_particle_count": 8000,
"star_population": {"G-Star": 1200, "BD": 800, "WD": 500, "NS": 300, "Magnetar": 10},
"NUM_TACHYONS": 5000,

# Output Control
"OUTPUT_FILENAME": "hilios_v18_final_state.h5",
"OUTPUT_DIR": "/content/drive/MyDrive/RDU_Sims/"
}

# =====
# === PART 2: CUDA KERNELS ===
# =====
merger_kernel = cp.RawKernel(r"""
extern "C" __global__
void merger_kernel(float* pos, float* vel, float* mass, float* radius, int* is_consumed, int
num_particles,
                int* merger_log, int* merger_log_count, int max_log_size) {
    int i = blockDim.x * blockIdx.x + threadIdx.x;
    if (i >= num_particles || is_consumed[i] == 1) return;

    for (int j = i + 1; j < num_particles; ++j) {
        if (is_consumed[j] == 1) continue;

        float dx = pos[i*3+0] - pos[j*3+0];
        float dy = pos[i*3+1] - pos[j*3+1];
        float dz = pos[i*3+2] - pos[j*3+2];
        float dist_sq = dx*dx + dy*dy + dz*dz;
        float combined_radius = radius[i] + radius[j];

        if (dist_sq < combined_radius * combined_radius) {
            int survivor_idx = (mass[i] >= mass[j]) ? i : j;
            int consumed_idx = (mass[i] >= mass[j]) ? j : i;

            if (atomicCAS(&is_consumed[consumed_idx], 0, 1) == 0) {
                float m_survivor = mass[survivor_idx];
                float m_consumed = mass[consumed_idx];
                float total_mass = m_survivor + m_consumed;

```

```

        vel[survivor_idx*3+0] = (m_survivor * vel[survivor_idx*3+0] + m_consumed *
vel[consumed_idx*3+0]) / total_mass;
        vel[survivor_idx*3+1] = (m_survivor * vel[survivor_idx*3+1] + m_consumed *
vel[consumed_idx*3+1]) / total_mass;
        vel[survivor_idx*3+2] = (m_survivor * vel[survivor_idx*3+2] + m_consumed *
vel[consumed_idx*3+2]) / total_mass;

        atomicAdd(&mass[survivor_idx], m_consumed);
        radius[survivor_idx] = cbrtf(powf(radius[survivor_idx], 3) + powf(radius[consumed_idx],
3));

        int log_idx = atomicAdd(merger_log_count, 1);
        if (log_idx < max_log_size) {
            merger_log[log_idx * 2 + 0] = survivor_idx;
            merger_log[log_idx * 2 + 1] = consumed_idx;
        }
    }
}
}
}
""", 'merger_kernel')

```

```

stellar_evolution_kernel = cp.RawKernel(r"""
extern "C" __global__
void stellar_evolution_kernel(int* particle_types, float* lifetimes, int* event_flags, int n, float dt) {
    int i = blockDim.x* blockIdx.x + threadIdx.x;
    if (i >= n || lifetimes[i] <= 0) return;
    lifetimes[i] -= dt;
    if (lifetimes[i] <= 0) {
        if (particle_types[i] == 2) {
            particle_types[i] = 4;
            event_flags[i] = 1;
        }
    }
}
}
""", 'stellar_evolution_kernel')

```

```

knoechelman_kernel = cp.RawKernel(r"""
extern "C" __global__
void knoechelman_kernel(const float* positions, float* velocities, float* masses, const float*
radii, const float* spins,
                        int* particle_types, int* is_consumed, int* event_flags, int n, float A, float N_sens,
float C_unify,
                        float k_rot, float C_balance, float G_rdu, float tachyon_mass) {

```

```

int i = blockDim.x * blockIdx.x + threadIdx.x;
if (i >= n || is_consumed[i] == 1 || particle_types[i] == 0 || particle_types[i] >= 10) return;

float drc_pos_x = positions[0], drc_pos_y = positions[1], drc_pos_z = positions[2];
float drc_radius = radii[0];
float dx = positions[i*3+0] - drc_pos_x;
float dy = positions[i*3+1] - drc_pos_y;
float dz = positions[i*3+2] - drc_pos_z;
float dist_sq = dx*dx + dy*dy + dz*dz;

if (dist_sq < (radii[i] + drc_radius) * (radii[i] + drc_radius)) {
    if (atomicExch(&is_consumed[i], 1) == 0) {
        float drc_mass = masses[0];
        float drc_spin_mag = sqrtf(spins[0]*spins[0] + spins[1]*spins[1] + spins[2]*spins[2]);
        float l_over_lmax_proxy = fminf(1.0f, drc_spin_mag / 100.0f);
        float drc_density = (drc_radius > 0) ? drc_mass / (4.1887f * drc_radius * drc_radius *
drc_radius) : 0.0f;
        float phi_U = C_unify * (drc_mass * drc_density) * expf(k_rot * l_over_lmax_proxy);
        float consumed_mass = masses[i];
        float omega_R = A * powf(phi_U / fmaxf(consumed_mass, 1e-9f), N_sens);
        float e_trap = (drc_radius > 0) ? G_rdu * (drc_mass * drc_mass) / drc_radius : 1e9f;
        float xi = (e_trap > 0) ? C_balance * (omega_R * omega_R) / e_trap : 0.0f;

        if (xi > 1.0f) {
            particle_types[i] = 11;
            velocities[i*3+0] *= 50.0f;
            velocities[i*3+1] *= 50.0f;
            velocities[i*3+2] *= 50.0f;
        } else {
            particle_types[i] = 10;
            velocities[i*3+0] *= 5.0f;
            velocities[i*3+1] *= 5.0f;
            velocities[i*3+2] *= 5.0f;
        }
        masses[i] = tachyon_mass;
        event_flags[i] = 2;
    }
}
}
}
""", 'knoechelman_kernel')

```

```

deposit_mass_kernel = cp.RawKernel(r"""
extern "C" __global__ void deposit_mass_kernel(const float* positions, const float* masses,
float* grid, int n, int gs) {

```

```

int i = blockDim.x * blockIdx.x + threadIdx.x;
if (i >= n) return;
int x = (int)positions[i*3+0], y = (int)positions[i*3+1], z = (int)positions[i*3+2];
if (x >= 0 && x < gs && y >= 0 && y < gs && z >= 0 && z < gs) {
    atomicAdd(&grid[x*gs*gs + y*gs + z], masses[i]);
}
}
"""', 'deposit_mass_kernel')

```

```

update_chronos_kernel = cp.RawKernel(r"""
extern "C" __global__ void update_chronos_kernel(float* C_n_plus_1, const float* C_n, const
float* C_n_minus_1, const float* source_term, int gs, float dt, float K_S, float G) {
    int x = blockIdx.x * blockDim.x + threadIdx.x, y = blockIdx.y * blockDim.y + threadIdx.y, z =
blockIdx.z * blockDim.z + threadIdx.z;
    if (x >= gs || y >= gs || z >= gs) return;
    if (x > 0 && x < gs - 1 && y > 0 && y < gs - 1 && z > 0 && z < gs - 1) {
        int idx = x*gs*gs + y*gs + z;
        float laplacian = K_S * (C_n[idx+gs*gs] + C_n[idx-gs*gs] + C_n[idx+gs] + C_n[idx-gs] +
C_n[idx+1] + C_n[idx-1] - 6.0f * C_n[idx]);
        float source = -G * source_term[idx];
        C_n_plus_1[idx] = 2.0f * C_n[idx] - C_n_minus_1[idx] + dt * dt * (laplacian + source);
    }
}
"""', 'update_chronos_kernel')

```

```

generate_b_field_kernel = cp.RawKernel(r"""
extern "C" __global__ void generate_b_field_kernel(float* b_field_grid, const float* positions,
const float* spins, const int* particle_types, int n, int gs, float b_strength) {
    int x = blockIdx.x * blockDim.x + threadIdx.x, y = blockIdx.y * blockDim.y + threadIdx.y, z =
blockIdx.z * blockDim.z + threadIdx.z;
    if (x >= gs || y >= gs || z >= gs) return;
    int grid_idx = (x*gs*gs + y*gs + z) * 3;
    float bx = 0.0f, by = 0.0f, bz = 0.0f;
    for (int i = 0; i < n; ++i) {
        if (particle_types[i] == 7) { // Magnetar
            float mx = positions[i*3+0], my = positions[i*3+1], mz = positions[i*3+2];
            float rx = (float)x - mx, ry = (float)y - my, rz = (float)z - mz;
            float dist_sq = rx*rx + ry*ry + rz*rz;
            if (dist_sq > 1e-6f) {
                float strength = b_strength / (dist_sq * sqrtf(dist_sq));
                float sx = spins[i*3+0], sy = spins[i*3+1], sz = spins[i*3+2];
                float smag = sqrtf(sx*sx + sy*sy + sz*sz);
                if (smag > 1e-6f) {
                    bx += strength * sx / smag;

```

```

        by += strength * sy / smag;
        bz += strength * sz / smag;
    }
}
}
}
b_field_grid[grid_idx+0] = bx;
b_field_grid[grid_idx+1] = by;
b_field_grid[grid_idx+2] = bz;
}
""", 'generate_b_field_kernel')

```

```

force_kernel = cp.RawKernel(r"""
extern "C" __global__
void force_kernel(
    float* accelerations, const float* positions, const float* velocities, const float* masses,
    const int* particle_types, const float* chronos_field, const float* b_field_grid,
    int n, int gs, float C_drag, float R_drag, float C_friction, float q,
    float dsf_strength, float dsf_radius) {

    int i = blockIdx.x * blockDim.x + threadIdx.x;
    if (i >= n) return;

    float px = positions[i*3+0], py = positions[i*3+1], pz = positions[i*3+2];
    float ax = 0.0f, ay = 0.0f, az = 0.0f;
    int x = (int)px, y = (int)py, z = (int)pz;

    // Force from Chronos Field Gradient
    if (x > 0 && x < gs - 1 && y > 0 && y < gs - 1 && z > 0 && z < gs - 1) {
        ax = -(chronos_field[(x+1)*gs*gs+y*gs+z] - chronos_field[(x-1)*gs*gs+y*gs+z]) / 2.0f;
        ay = -(chronos_field[x*gs*gs+(y+1)*gs+z] - chronos_field[x*gs*gs+(y-1)*gs+z]) / 2.0f;
        az = -(chronos_field[x*gs*gs+y*gs+(z+1)] - chronos_field[x*gs*gs+y*gs+(z-1)]) / 2.0f;
    }

    int p_type = particle_types[i];

    // Gas-specific forces
    if (p_type == 1) {
        // Lorentz force for charged gas
        if (x >= 0 && x < gs && y >= 0 && y < gs && z >= 0 && z < gs) {
            int b_idx = (x*gs*gs + y*gs + z) * 3;
            float bx = b_field_grid[b_idx+0], by = b_field_grid[b_idx+1], bz = b_field_grid[b_idx+2];
            float vx = velocities[i*3+0], vy = velocities[i*3+1], vz = velocities[i*3+2];
            ax += q * (vy * bz - vz * by);

```

```

    ay += q * (vz * bx - vx * bz);
    az += q * (vx * by - vy * bx);
}
// Accretion drag near the center
float d_sq_c = (50.5f - px)*(50.5f - px) + (50.5f - py)*(50.5f - py) + (50.5f - pz)*(50.5f - pz);
if (d_sq_c < R_drag * R_drag) {
    ax -= C_drag * velocities[i*3+0];
    ay -= C_drag * velocities[i*3+1];
    az -= C_drag * velocities[i*3+2];
}
}

// Dynamical Friction for massive objects
if (masses[i] > 10.0) {
    float v_sq = velocities[i*3+0]*velocities[i*3+0] + velocities[i*3+1]*velocities[i*3+1] +
    velocities[i*3+2]*velocities[i*3+2];
    if (v_sq > 0.1f) {
        float f_mag = C_friction * masses[i] / v_sq;
        ax -= f_mag * velocities[i*3+0];
        ay -= f_mag * velocities[i*3+1];
        az -= f_mag * velocities[i*3+2];
    }
}

// Dark Strong Force between Tachyons
if (p_type >= 10) {
    for (int j = 0; j < n; ++j) {
        if (i == j || particle_types[j] < 10) continue;
        float dx = px - positions[j*3+0], dy = py - positions[j*3+1], dz = pz - positions[j*3+2];
        float r_sq = dx*dx + dy*dy + dz*dz;
        if (r_sq > 1e-6f && r_sq < 25.0f) {
            float r = sqrtf(r_sq);
            float r0_over_r_6 = powf(dsf_radius / r, 6.0f);
            float r0_over_r_12 = r0_over_r_6 * r0_over_r_6;
            float force_mag = (12.0f * dsf_strength / r) * (r0_over_r_12 - r0_over_r_6);
            ax += force_mag * (dx / r) / masses[i];
            ay += force_mag * (dy / r) / masses[i];
            az += force_mag * (dz / r) / masses[i];
        }
    }
}

accelerations[i*3+0] = ax;
accelerations[i*3+1] = ay;

```

```

    accelerations[i*3+2] = az;
}
"""', 'force_kernel')

# =====
# === PART 3: DATA INITIALIZATION & GPU TRANSFER ===
# =====
print("Initializing Hilios v18...");
star_counts = sim_params["star_population"];
num_stars = sum(star_counts.values())
num_baryons = sim_params["gas_particle_count"] + num_stars + 1; # +1 for DRC
num_tachyons = sim_params["NUM_TACHYONS"]
total_particles = num_baryons + num_tachyons;
gs = sim_params["GRID_SIZE"]

positions_cpu = np.zeros((total_particles, 3), dtype=np.float32);
velocities_cpu = np.zeros((total_particles, 3), dtype=np.float32)
masses_cpu = np.ones(total_particles, dtype=np.float32);
radii_cpu = np.ones(total_particles, dtype=np.float32) * 0.5
is_consumed_cpu = np.zeros(total_particles, dtype=np.int32);
spin_cpu = np.zeros((total_particles, 3), dtype=np.float32)
particle_type_cpu = np.zeros(total_particles, dtype=np.int32);
lifetimes_cpu = np.zeros(total_particles, dtype=np.float32)

print(f'Creating {total_particles} particles...');
# Dark Resonator Core (DRC)
particle_type_cpu[0] = 0;
masses_cpu[0] = 50000.0;
radii_cpu[0] = 2.0
spin_cpu[0] = [0, 0, 100.0];
positions_cpu[0] = [gs / 2, gs / 2, gs / 2];

# Accretion Disk
disk_indices = slice(1, num_baryons)
num_disk_particles = num_baryons - 1;
radius = np.random.uniform(gs * 0.1, gs * 0.5, num_disk_particles);
angle = np.random.uniform(0, 2 * np.pi, num_disk_particles);
z_pos = np.random.uniform(-gs * 0.04, gs * 0.04, num_disk_particles)
positions_cpu[disk_indices, 0] = gs / 2 + radius * np.cos(angle);
positions_cpu[disk_indices, 1] = gs / 2 + radius * np.sin(angle)
positions_cpu[disk_indices, 2] = gs / 2 + z_pos;

orbital_v = np.sqrt(sim_params["G_RDU"] * masses_cpu[0] / radius)
velocities_cpu[disk_indices, 0] = -orbital_v * np.sin(angle);

```



```

velocities_cpu[disk_indices, 1] = orbital_v * np.cos(angle)

# Populate particle types in the disk
start_idx = 1;
particle_type_cpu[start_idx:start_idx + sim_params["gas_particle_count"]] = 1; # Gas
start_idx += sim_params["gas_particle_count"]

type_map = {"G-Star": 2, "BD": 3, "WD": 4, "NS": 5, "Magnetar": 7};
label_map = {v: k for k, v in type_map.items()}
label_map[0] = "DRC";
label_map[1] = "Gas"

for star_type, count in star_counts.items():
    end_idx = start_idx + count
    particle_type_cpu[start_idx:end_idx] = type_map[star_type]
    if star_type == "NS":
        lifetimes_cpu[start_idx:end_idx] = np.random.uniform(0.1, sim_params["NUM_STEPS"] *
sim_params["DT"])
    start_idx = end_idx

# --- NEW: Assign physically-based masses to stars ---
print("Assigning star masses...");
mass_map = {
    2: 1.0, # G-Star (Solar mass)
    3: 0.06, # BD (Brown Dwarf, e.g., 60 Jupiter masses)
    4: 0.9, # WD (White Dwarf)
    5: 1.8, # NS (Neutron Star)
    7: 2.0 # Magnetar
}
# Loop through all baryonic particles to set star masses
for i in range(1, num_baryons):
    p_type = particle_type_cpu[i]
    if p_type in mass_map:
        # Assign the mapped mass, with a small random variation
        masses_cpu[i] = mass_map[p_type] + np.random.uniform(-0.05, 0.05)
# Ensure gas particles remain at mass 1.0
gas_mask = (particle_type_cpu == 1)
masses_cpu[gas_mask] = 1.0

# Tachyons
tachyon_indices = slice(num_baryons, total_particles);
positions_cpu[tachyon_indices] = np.random.rand(num_tachyons, 3) * gs
velocities_cpu[tachyon_indices] = (np.random.rand(num_tachyons, 3) - 0.5) * 100;

```

```
particle_type_cpu[tachyon_indices] = 10
masses_cpu[tachyon_indices] = sim_params["TACHYON_MASS"] # Also assign tachyon mass
```

```
print("Transferring data to GPU...");
positions_gpu = cp.asarray(positions_cpu);
velocities_gpu = cp.asarray(velocities_cpu)
masses_gpu = cp.asarray(masses_cpu);
radii_gpu = cp.asarray(radii_cpu);
is_consumed_gpu = cp.asarray(is_consumed_cpu)
spin_gpu = cp.asarray(spin_cpu);
particle_type_gpu = cp.asarray(particle_type_cpu);
lifetimes_gpu = cp.asarray(lifetimes_cpu)
```

```
# Initialize GPU arrays
accelerations_gpu = cp.empty_like(positions_gpu);
event_flags_gpu = cp.zeros(total_particles, dtype=cp.int32)
chronos_field_n_minus_1 = cp.zeros((gs, gs, gs), dtype=cp.float32);
chronos_field_n = cp.zeros((gs, gs, gs), dtype=cp.float32)
chronos_field_n_plus_1 = cp.zeros((gs, gs, gs), dtype=cp.float32);
source_term_gpu = cp.zeros((gs, gs, gs), dtype=cp.float32)
b_field_grid = cp.zeros((gs, gs, gs, 3), dtype=cp.float32)
```

```
# Merger logging setup
MAX_MERGERS_PER_STEP = 100;
merger_log_gpu = cp.zeros((MAX_MERGERS_PER_STEP, 2), dtype=cp.int32);
merger_log_count_gpu = cp.zeros(1, dtype=cp.int32)
```

```
print("Initialization complete.")
```

```
# =====
# === PART 4: LIVE SCIENCE DASHBOARD ===
# =====
fig = plt.figure(figsize=(16, 12));
fig.patch.set_facecolor('#0d0d0d');
dashboard_gs = gridspec.GridSpec(2, 4, height_ratios=[3, 1.2])
ax_main = fig.add_subplot(dashboard_gs[0, :], projection='3d');
ax_hist_baryon = fig.add_subplot(dashboard_gs[1, 0])
ax_hist_tachyon = fig.add_subplot(dashboard_gs[1, 1]);
ax_angular_momentum = fig.add_subplot(dashboard_gs[1, 2])
ax_log = fig.add_subplot(dashboard_gs[1, 3]);
plt.subplots_adjust(wspace=0.3, hspace=0.5)
```

```
# UPDATED Dashboard Function
```

```

def update_dashboard(step, positions, velocities, particle_types, masses, ang_mom_history,
event_log_list):
    # --- Clear all axes for redrawing ---
    for ax in [ax_main, ax_hist_baryon, ax_hist_tachyon, ax_angular_momentum, ax_log]:
        ax.cla()

    # --- Main 3D Plot ---
    ax_main.set_facecolor('black')
    ax_main.set_title(f'Galaxy State at Step {step}', color='white', fontsize=16)

    type_color_map = {0: "magenta", 1: "#333333", 2: "yellow", 3: "orange", 4: "white", 5: "cyan",
7: "lime", 10: "red", 11: "purple"}
    type_size_map = {0: 200, 1: 2, 2: 15, 3: 10, 4: 12, 5: 20, 7: 25, 10: 5, 11: 5}

    unique_types = np.unique(particle_types)
    for p_type in unique_types:
        mask = (particle_types == p_type)
        marker = "*" if p_type == 0 else "o"
        alpha = 0.1 if p_type == 1 or p_type >= 10 else 1.0
        ax_main.scatter(positions[mask, 0], positions[mask, 1], positions[mask, 2],
c=type_color_map.get(p_type), s=type_size_map.get(p_type), alpha=alpha, marker=marker)

    plot_limit = sim_params["GRID_SIZE"]
    ax_main.set(xlim=(0, plot_limit), ylim=(0, plot_limit), zlim=(0, plot_limit))
    for axis in [ax_main.xaxis, ax_main.yaxis, ax_main.zaxis]:
        axis.pane.fill = False
        axis.line.set_color('grey')
        axis.set_tick_params(colors='grey')

    # --- Baryon Mass Histogram ---
    baryon_mask = (particle_types < 10) & (particle_types > 0)
    ax_hist_baryon.hist(masses[baryon_mask], bins=50, color='cyan', alpha=0.7) # Increased
bins
    ax_hist_baryon.set_title('Baryon Mass Dist.', color='white')
    ax_hist_baryon.set_facecolor('#1a1a1a')
    ax_hist_baryon.tick_params(colors='grey')

    # --- Tachyon Velocity Histogram ---
    tachyon_mask = (particle_types >= 10)
    if np.any(tachyon_mask):
        tachyon_speeds = np.linalg.norm(velocities[tachyon_mask], axis=1)
        ax_hist_tachyon.hist(tachyon_speeds, bins=20, color='red', alpha=0.7)
    ax_hist_tachyon.set_title('Tachyon Speed Dist.', color='white')
    ax_hist_tachyon.set_facecolor('#1a1a1a')

```

```

ax_hist_tachyon.tick_params(colors='grey')

# --- Angular Momentum Plot ---
steps_plotted = range(0, len(ang_mom_history) * sim_params["VIS_INTERVAL"],
sim_params["VIS_INTERVAL"])
if len(steps_plotted) > 1:
    ax_angular_momentum.plot(steps_plotted, ang_mom_history, color='lime')
    ax_angular_momentum.set_title('Ang. Momentum', color='white')
    ax_angular_momentum.set_facecolor('#1a1a1a')
    ax_angular_momentum.tick_params(colors='grey')

# --- Event Log ---
ax_log.set_title('Event Log', color='white')
ax_log.set_facecolor('#1a1a1a')
ax_log.set_xticks([])
ax_log.set_yticks([])
ax_log.text(0.05, 0.95, '\n'.join(event_log_list[-5:]), transform=ax_log.transAxes,
color='lightgreen', fontsize=10, va='top', family='monospace')

# --- Display Update ---
display(fig)
clear_output(wait=True)

# =====
# === PART 5: DATA SAVING FUNCTION ===
# =====
def save_final_state(filepath, pos_gpu, vel_gpu, types_gpu, masses_gpu, spin_gpu):
    print(f"\nSaving final simulation state to: {filepath}");
    try:
        os.makedirs(os.path.dirname(filepath), exist_ok=True)
        pos_host = cp.asnumpy(pos_gpu);
        vel_host = cp.asnumpy(vel_gpu);
        types_host = cp.asnumpy(types_gpu)
        masses_host = cp.asnumpy(masses_gpu);
        spin_host = cp.asnumpy(spin_gpu)
        with h5py.File(filepath, 'w') as f:
            p_group = f.create_group('particles')
            p_group.create_dataset('positions', data=pos_host);
            p_group.create_dataset('velocities', data=vel_host)
            p_group.create_dataset('types', data=types_host);
            p_group.create_dataset('masses', data=masses_host)
            p_group.create_dataset('spins', data=spin_host)
        print("Save complete.")
    except Exception as e:

```

```

print(f"--- ERROR: Could not save file. ---\n{e}")

# =====
# === PART 6: MAIN SIMULATION LOOP ===
# =====
print(f"Starting Hilios v18 with {total_particles} particles...")
threads_per_block = 256;
blocks_per_grid_1d = (total_particles + threads_per_block - 1) // threads_per_block
event_log = ["Sim Initialized..."]
angular_momentum_history = []

for step in range(sim_params["NUM_STEPS"]):
    # --- Merger Detection ---
    merger_log_count_gpu.fill(0)
    merger_kernel((blocks_per_grid_1d,), (threads_per_block,),
                  (positions_gpu, velocities_gpu, masses_gpu, radii_gpu, is_consumed_gpu,
                   total_particles, merger_log_gpu, merger_log_count_gpu, MAX_MERGERS_PER_STEP))
    num_mergers = int(merger_log_count_gpu.get().item())

    # FIX: Cap the number of mergers to the maximum we can log to prevent crashes.
    num_mergers = min(num_mergers, MAX_MERGERS_PER_STEP)

    if num_mergers > 0:
        merger_events = cp.asnumpy(merger_log_gpu[:num_mergers]);
        involved_indices = merger_events.flatten()
        involved_types = cp.asnumpy(particle_type_gpu[involved_indices]).reshape(num_mergers,
2)
        for i in range(num_mergers):
            survivor_name = label_map.get(involved_types[i, 0], "Unk");
            consumed_name = label_map.get(involved_types[i, 1], "Unk")
            event_log.append(f"MERGER @ {step}: {survivor_name} + {consumed_name}")

    # --- Chronos Field Update ---
    source_term_gpu.fill(0)
    deposit_mass_kernel((blocks_per_grid_1d,), (threads_per_block,),
                        (positions_gpu, masses_gpu, source_term_gpu, total_particles, gs))

    threads_per_block_3d = (8, 8, 8);
    blocks_per_grid_3d = ((gs + 7) // 8, (gs + 7) // 8, (gs + 7) // 8)

    # FIX: Corrected syntax for 3D kernel calls
    update_chronos_kernel(blocks_per_grid_3d, threads_per_block_3d,
                          (chronos_field_n_plus_1, chronos_field_n, chronos_field_n_minus_1,
                           source_term_gpu, gs, sim_params["DT"], sim_params["K_S"], sim_params["G_RDU"]))

```

```

chronos_field_n, chronos_field_n_minus_1 = chronos_field_n_plus_1, chronos_field_n

# --- B-Field Generation ---
b_field_grid.fill(0)
# FIX: Corrected syntax for 3D kernel calls
generate_b_field_kernel(blocks_per_grid_3d, threads_per_block_3d,
                        (b_field_grid, positions_gpu, spin_gpu, particle_type_gpu, total_particles, gs,
sim_params["MAGNETAR_STRENGTH"])))

# --- Force Calculation & Integration ---
force_kernel((blocks_per_grid_1d,), (threads_per_block,),
            (accelerations_gpu, positions_gpu, velocities_gpu, masses_gpu, particle_type_gpu,
chronos_field_n.ravel(), b_field_grid, total_particles, gs, sim_params["C_ACCRETION_DRAG"],
sim_params["ACCRETION_RADIUS"], sim_params["C_DYN_FRICTION"],
sim_params["GAS_CHARGE"], sim_params["DARK_FORCE_STRENGTH"],
sim_params["REPULSION_RADIUS"])))

velocities_gpu += accelerations_gpu * sim_params["DT"];
positions_gpu += velocities_gpu * sim_params["DT"]

# --- Stellar Evolution & Supernovae ---
event_flags_gpu.fill(0)
stellar_evolution_kernel((blocks_per_grid_1d,), (threads_per_block,),
                        (particle_type_gpu, lifetimes_gpu, event_flags_gpu, total_particles,
sim_params["DT"])))

if cp.any(event_flags_gpu == 1):
    supernova_indices_gpu = cp.where(event_flags_gpu == 1)[0]
    if supernova_indices_gpu.size > 0:
        log_message = f"SUPERNOVA @ {step}: {supernova_indices_gpu.size} star(s)
exploded!";
        event_log.append(log_message)
        supernova_positions = cp.asnumpy(positions_gpu[supernova_indices_gpu]);
        supernova_masses = cp.asnumpy(masses_gpu[supernova_indices_gpu])
        new_particles_pos, new_particles_vel, new_particles_mass, new_particles_type,
new_particles_radius, new_particles_lifetimes, new_particles_spins = [], [], [], [], [], []

        for i in range(supernova_indices_gpu.size):
            pos, original_mass = supernova_positions[i], supernova_masses[i]
            remnant_mass = original_mass * sim_params["REMNANT_MASS_FRACTION"]
            new_particles_pos.append(pos);
            new_particles_vel.append([0, 0, 0]);
            new_particles_mass.append(remnant_mass)

```

```

if original_mass > 2.5:
    new_particles_type.append(7) # Magnetar
else:
    new_particles_type.append(5) # Neutron Star

new_particles_radius.append(0.01);
new_particles_lifetimes.append(1e18);
new_particles_spins.append([0, 0, 0])

shell_particles = sim_params["SUPERNOVA_SHELL_PARTICLES"];
ejection_vel = sim_params["SUPERNOVA_EJECTION_VELOCITY"]
mass_per_shell_particle = (original_mass * (1.0 -
sim_params["REMNANT_MASS_FRACTION"])) / shell_particles

for _ in range(shell_particles):
    theta, phi = np.arccos(2 * np.random.rand() - 1), np.random.uniform(0, 2 * np.pi)
    vx, vy, vz = ejection_vel * np.sin(theta) * np.cos(phi), ejection_vel * np.sin(theta) *
np.sin(phi), ejection_vel * np.cos(theta)
    new_particles_pos.append(pos);
    new_particles_vel.append([vx, vy, vz]);
    new_particles_mass.append(mass_per_shell_particle)
    new_particles_type.append(1); # Gas
    new_particles_radius.append(0.001);
    new_particles_lifetimes.append(1e18);
    new_particles_spins.append([0, 0, 0])

# Remove original stars and add new remnants and shell particles
keep_mask = (event_flags_gpu != 1)
positions_gpu, velocities_gpu, masses_gpu, radii_gpu, is_consumed_gpu, spin_gpu,
particle_type_gpu, lifetimes_gpu = \
    positions_gpu[keep_mask], velocities_gpu[keep_mask], masses_gpu[keep_mask],
radii_gpu[keep_mask], is_consumed_gpu[keep_mask], spin_gpu[keep_mask],
particle_type_gpu[keep_mask], lifetimes_gpu[keep_mask]

num_new_particles = len(new_particles_pos)
new_pos_gpu = cp.asarray(np.array(new_particles_pos));
new_vel_gpu = cp.asarray(np.array(new_particles_vel));
new_mass_gpu = cp.asarray(np.array(new_particles_mass))
new_type_gpu = cp.asarray(np.array(new_particles_type), dtype=cp.int32);
new_radius_gpu = cp.asarray(np.array(new_particles_radius));
new_lifetimes_gpu = cp.asarray(np.array(new_particles_lifetimes));
new_spins_gpu = cp.asarray(np.array(new_spins_gpu))
new_consumed_gpu = cp.zeros(num_new_particles, dtype=cp.int32)

```

```

positions_gpu = cp.vstack([positions_gpu, new_pos_gpu]);
velocities_gpu = cp.vstack([velocities_gpu, new_vel_gpu]);
masses_gpu = cp.hstack([masses_gpu, new_mass_gpu])
radii_gpu = cp.hstack([radii_gpu, new_radius_gpu]);
is_consumed_gpu = cp.hstack([is_consumed_gpu, new_consumed_gpu]);
spin_gpu = cp.vstack([spin_gpu, new_spins_gpu])
particle_type_gpu = cp.hstack([particle_type_gpu, new_type_gpu]);
lifetimes_gpu = cp.hstack([lifetimes_gpu, new_lifetimes_gpu])

# Update particle count and kernel launch parameters
total_particles = positions_gpu.shape[0];
event_flags_gpu = cp.zeros(total_particles, dtype=cp.int32)
blocks_per_grid_1d = (total_particles + threads_per_block - 1) // threads_per_block

# --- Knoechelman Mechanism ---
knoechelman_kernel((blocks_per_grid_1d,), (threads_per_block,),
                    (positions_gpu, velocities_gpu, masses_gpu, radii_gpu, spin_gpu,
particle_type_gpu, is_consumed_gpu, event_flags_gpu, total_particles,
                    sim_params["A_CONST"], sim_params["N_SENSITIVITY"],
sim_params["C_UNIFY"], sim_params["K_ROT"], sim_params["C_BALANCE"],
sim_params["G_RDU"], sim_params["TACHYON_MASS"])))

if cp.any(event_flags_gpu == 2):
    num_events = int(cp.sum(event_flags_gpu == 2).get().item());
    event_log.append(f"TRANSFORMATION @ {step}: {num_events} particle(s) converted.")

# --- Visualization ---
if step > 0 and step % sim_params["VIS_INTERVAL"] == 0:
    active_mask = (is_consumed_gpu == 0)
    pos_host = cp.asnumpy(positions_gpu[active_mask]);
    vel_host = cp.asnumpy(velocities_gpu[active_mask]);
    types_host = cp.asnumpy(particle_type_gpu[active_mask])
    masses_host = cp.asnumpy(masses_gpu[active_mask])

# Calculate Angular Momentum
if pos_host.shape[0] > 1:
    # Center positions relative to the DRC (particle 0)
    rel_pos = pos_host[1:] - pos_host[0]
    # Get masses for non-DRC particles
    rel_masses = masses_host[1:]
    # Get velocities for non-DRC particles
    rel_vel = vel_host[1:]
    # Calculate angular momentum for each particle ( $L = r \times p = m * (r \times v)$ )

```



```

    L_vectors = rel_masses[:, np.newaxis] * np.cross(rel_pos, rel_vel)
    # Total angular momentum is the sum of individual vectors
    L_total_vector = np.sum(L_vectors, axis=0)
    # Magnitude of the total angular momentum
    L_total_mag = np.linalg.norm(L_total_vector)
    angular_momentum_history.append(L_total_mag)
else:
    angular_momentum_history.append(0)

# UPDATED Function Call
update_dashboard(step, pos_host, vel_host, types_host, masses_host,
angular_momentum_history, event_log)

print("Simulation Finished.")

# =====
# === PART 7: FINAL SAVE ===
# =====
save_final_state(os.path.join(sim_params["OUTPUT_DIR"],
sim_params["OUTPUT_FILENAME"]),
                positions_gpu, velocities_gpu, particle_type_gpu, masses_gpu, spin_gpu)

```